

Convolutional Neural Networks

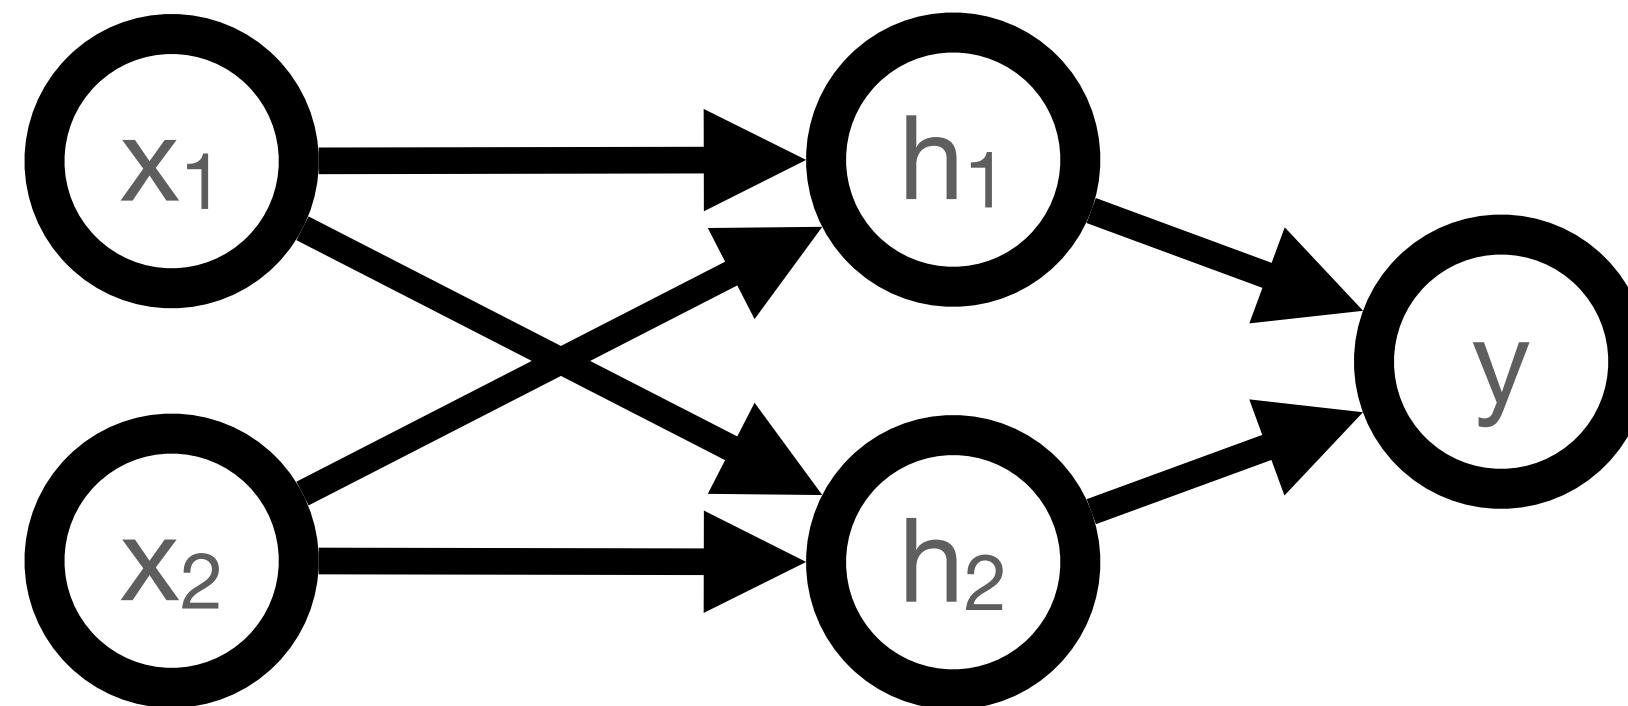
CMPUT 366: Intelligent Systems

GBC §9.0-9.4

Logistics

- Midterm is next **Monday, March 15**
 - Exam delivered via eclass
 - 24 hour window to take the exam
- There will be a review class this Friday
 - More details about format, likely questions, etc.

Recap: Feedforward Neural Network



- A **neural network** is many **units** **composed** together
- **Feedforward neural network:** Units arranged into **layers**
 - Each layer takes outputs of **previous layer** as its **inputs**

$$\begin{aligned} h_1(\mathbf{x}; \mathbf{w}^{(1)}, b^{(1)}) &= g \left(b^{(1)} + \sum_{i=1}^n w_i^{(1)} x_i \right) \\ y(\mathbf{x}; \mathbf{w}, \mathbf{b}) &= g \left(b^{(y)} + \sum_{i=1}^n w_i^{(y)} h_i(\mathbf{x}_i; \mathbf{w}^{(i)}, b^{(i)}) \right) \\ &= g \left(b^{(y)} + \sum_{i=1}^n w_i^{(y)} g \left(b^{(i)} + \sum_{j=1}^n w_j^{(i)} x_j \right) \right) \end{aligned}$$

Recap: Training Neural Networks

- Specify a **loss** L and a set of **training examples**:

$$E = (\mathbf{x}^{(1)}, y^{(1)}), \dots, (\mathbf{x}^{(n)}, y^{(n)})$$

- Training by **gradient descent**:

1. Compute **loss** on training data: $L(\mathbf{W}, \mathbf{b}) = \sum_i \ell(f(\mathbf{x}^{(i)}; \mathbf{W}, \mathbf{b}), y^{(i)})$

Loss function
(e.g., squared error)

i Prediction Target

2. Compute **gradient** of loss: $\nabla L(\mathbf{W}, \mathbf{b})$

3. **Update parameters** to make loss smaller:

$$\begin{bmatrix} \mathbf{W}^{new} \\ \mathbf{b}^{new} \end{bmatrix} = \begin{bmatrix} \mathbf{W}^{old} \\ \mathbf{b}^{old} \end{bmatrix} - \eta \nabla L(\mathbf{W}^{old}, \mathbf{b}^{old})$$

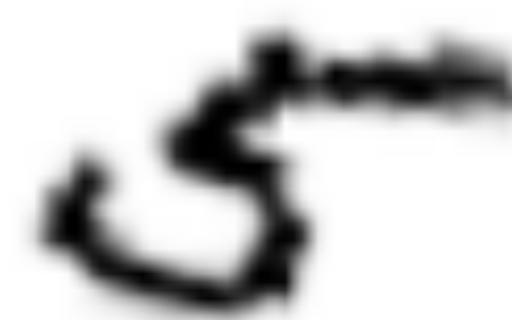
Recap: Automatic Differentiation

- **Forward mode** sweeps through the graph, computing $s'_i = \frac{\partial s_i}{\partial s_1}$ for each s_i
 - The **numerator varies**, and the **denominator is fixed**
 - At the end, we have computed $s'_n = \frac{\partial s_n}{\partial x_i}$ for a **single** input x_i
- **Backward mode** does the opposite:
 - For each s_i , computes the **local gradient** $\bar{s}_i = \frac{\partial s_n}{\partial s_i}$
 - The **numerator is fixed**, and the **denominator varies**
 - At the end, we have computed $\bar{x}_i = \frac{\partial s_n}{\partial x_i}$ for each input x_i
- **Key point:** The intermediate results are computed **numerically** at **each step**

Lecture Outline

1. Recap
2. Neural Networks for Image Recognition
3. Convolutional Neural Networks

Image Classification



FIVE

Problem: Recognize the handwritten digit from an image

- What are the **inputs**?
- What are the **outputs**?
- What is the **loss**?

Image Classification with Neural Networks

How can we use a **neural network** to solve this problem?

- How to represent the **inputs**?
- How to represent the **outputs**?
- What are the **parameters**?
- What is the **loss**?

1

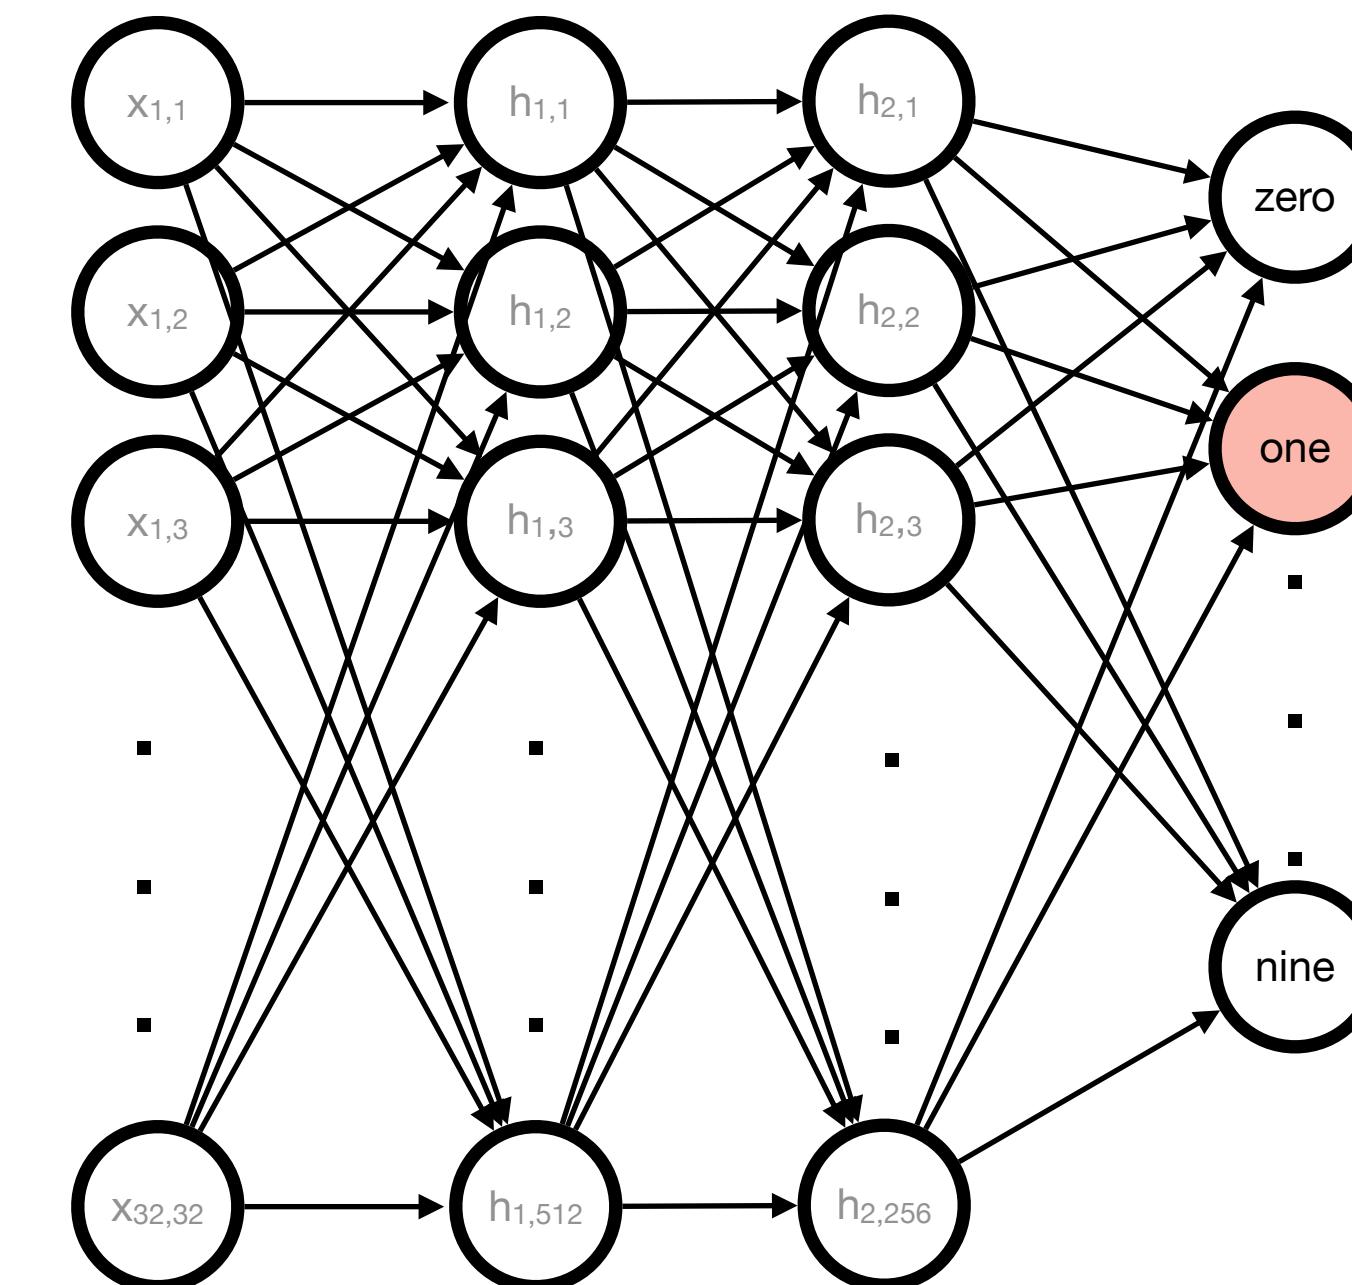
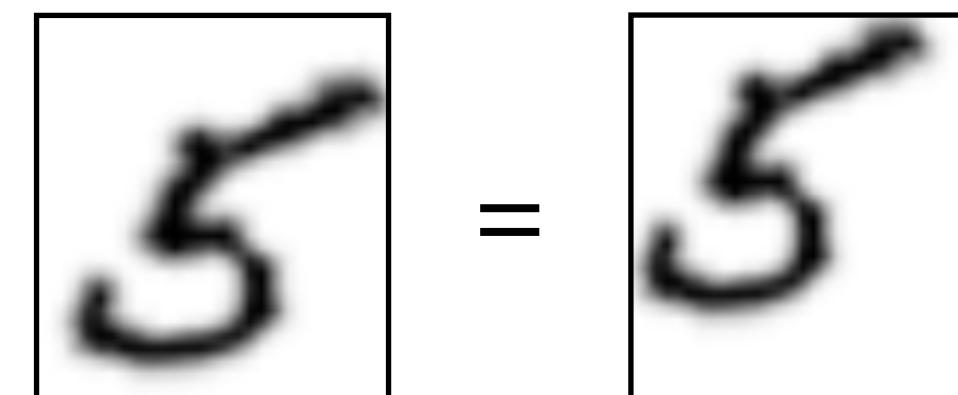
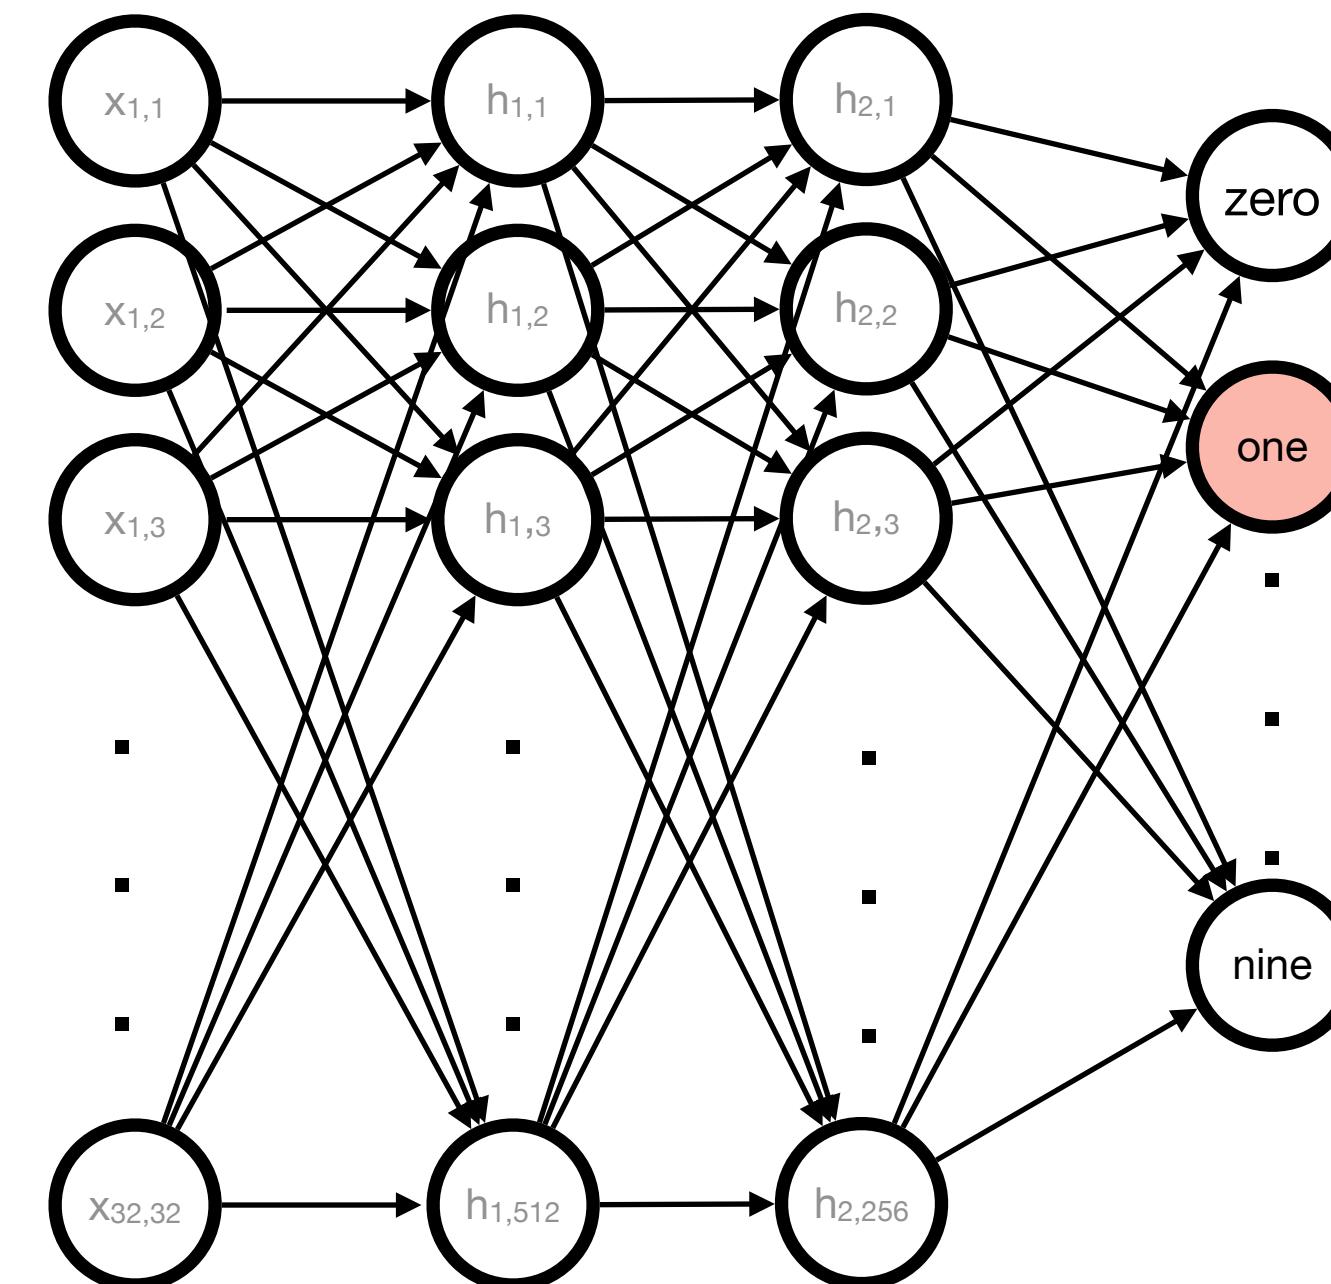


Image Recognition Issues

- For a large image, the number of parameters will be **very large**
 - For 32x32 greyscale image, hidden layer of 512 units hidden layer of 256 units, $1024 \times 512 + 512 \times 256 + 256 \times 10 = 657,920$ weights (and 1802 offsets)
 - Needs **lots of data** to train
 - Want to **generalize** over **transformations** of the input

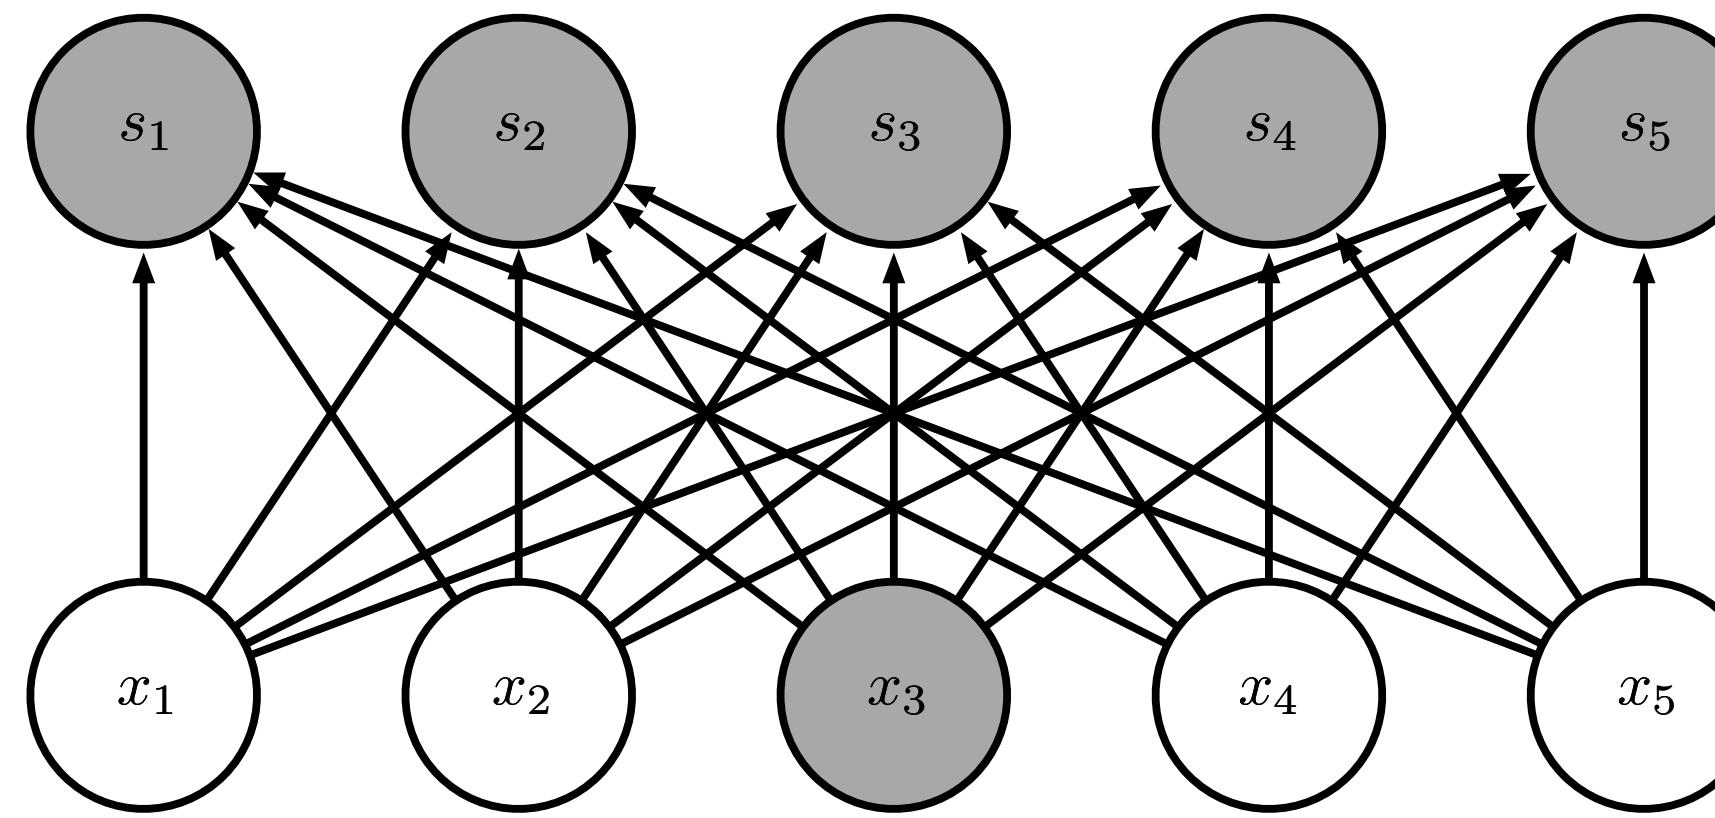


Convolutional Neural Networks

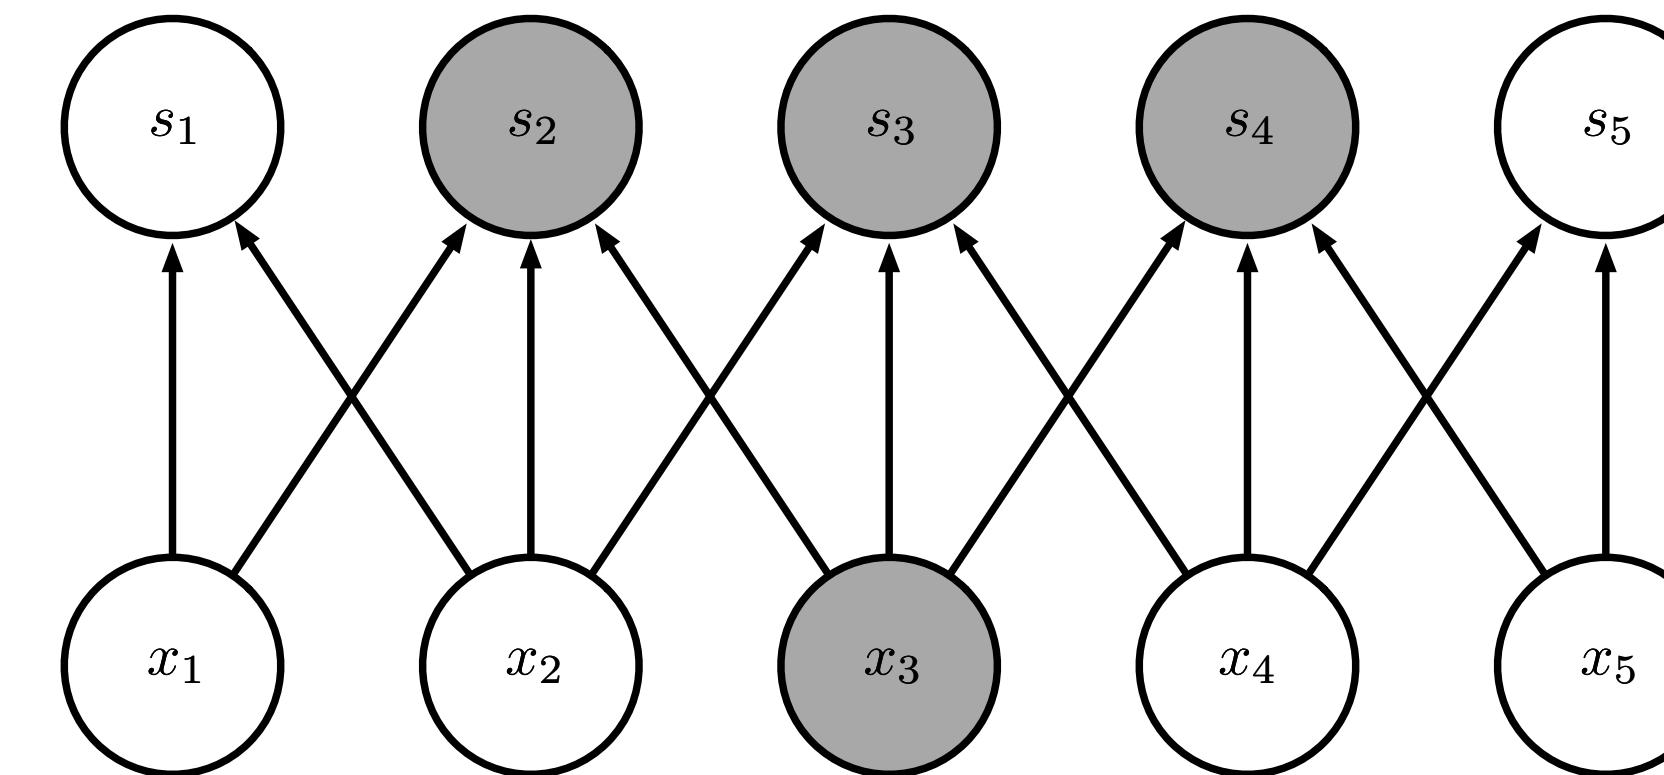
- **Convolutional neural networks:** a **specialized architecture** for image recognition
- Introduce two **new operations:**
 1. Convolutions
 2. Pooling
- Efficient **learning** via:
 1. Sparse interactions
 2. Parameter sharing
 3. Equivariant representations

Sparse Interactions

Dense
connections



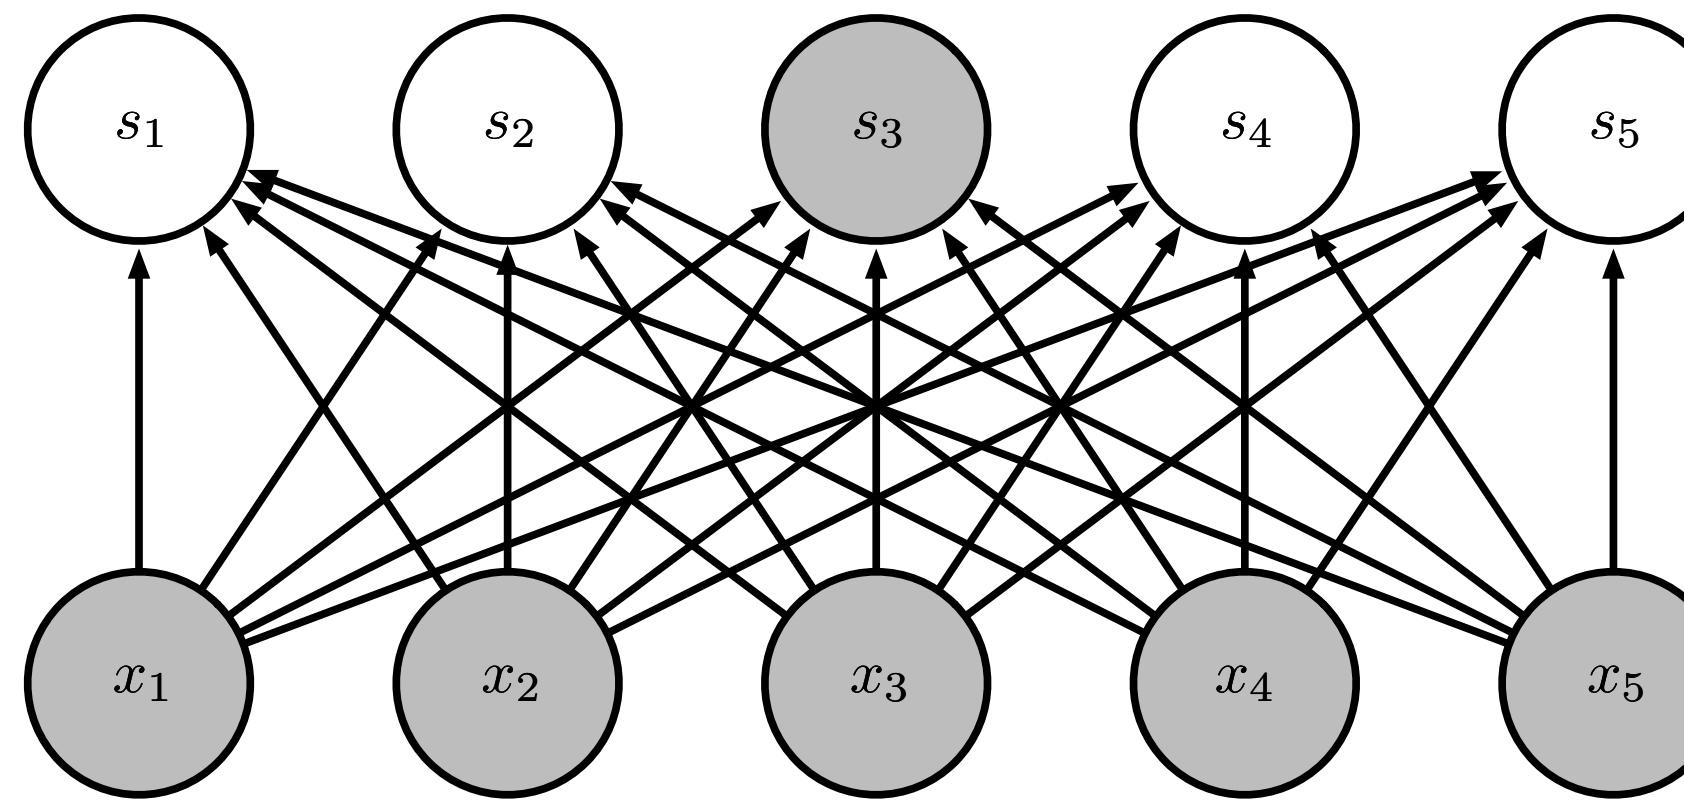
Sparse
connections



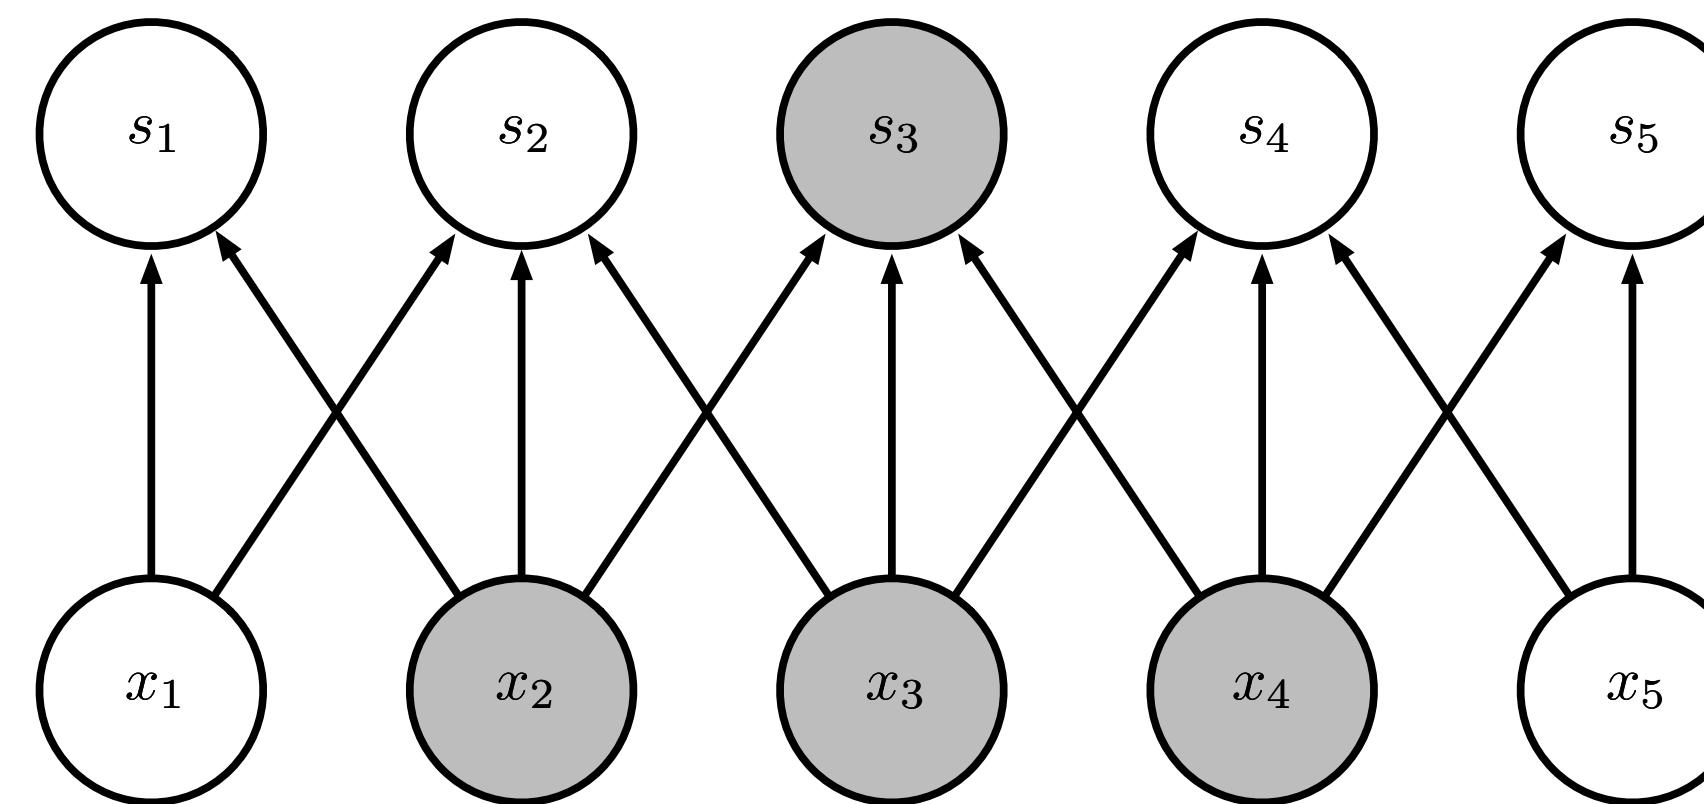
(Images: Goodfellow 2016)

Sparse Interactions

Dense
connections



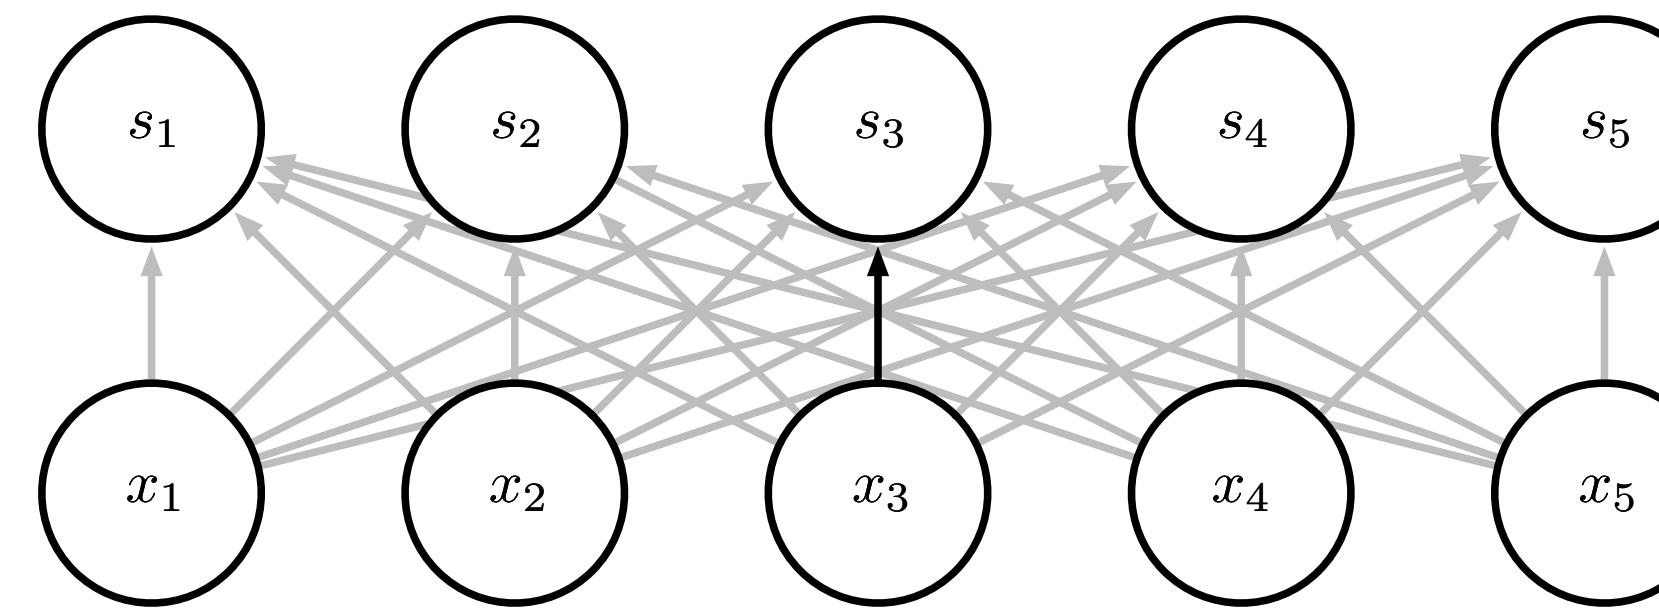
Sparse
connections



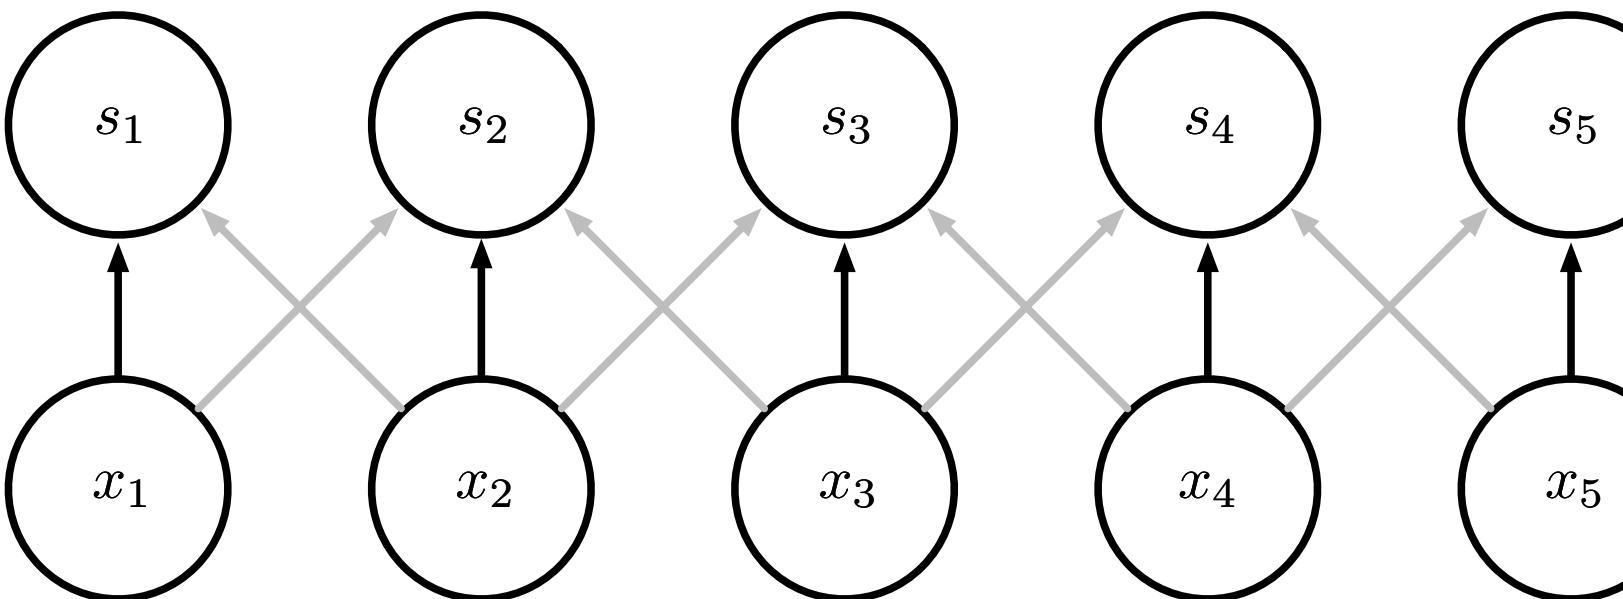
(Images: Goodfellow 2016)

Parameter Sharing

Traditional neural nets
learn a **unique value**
for **each connection**



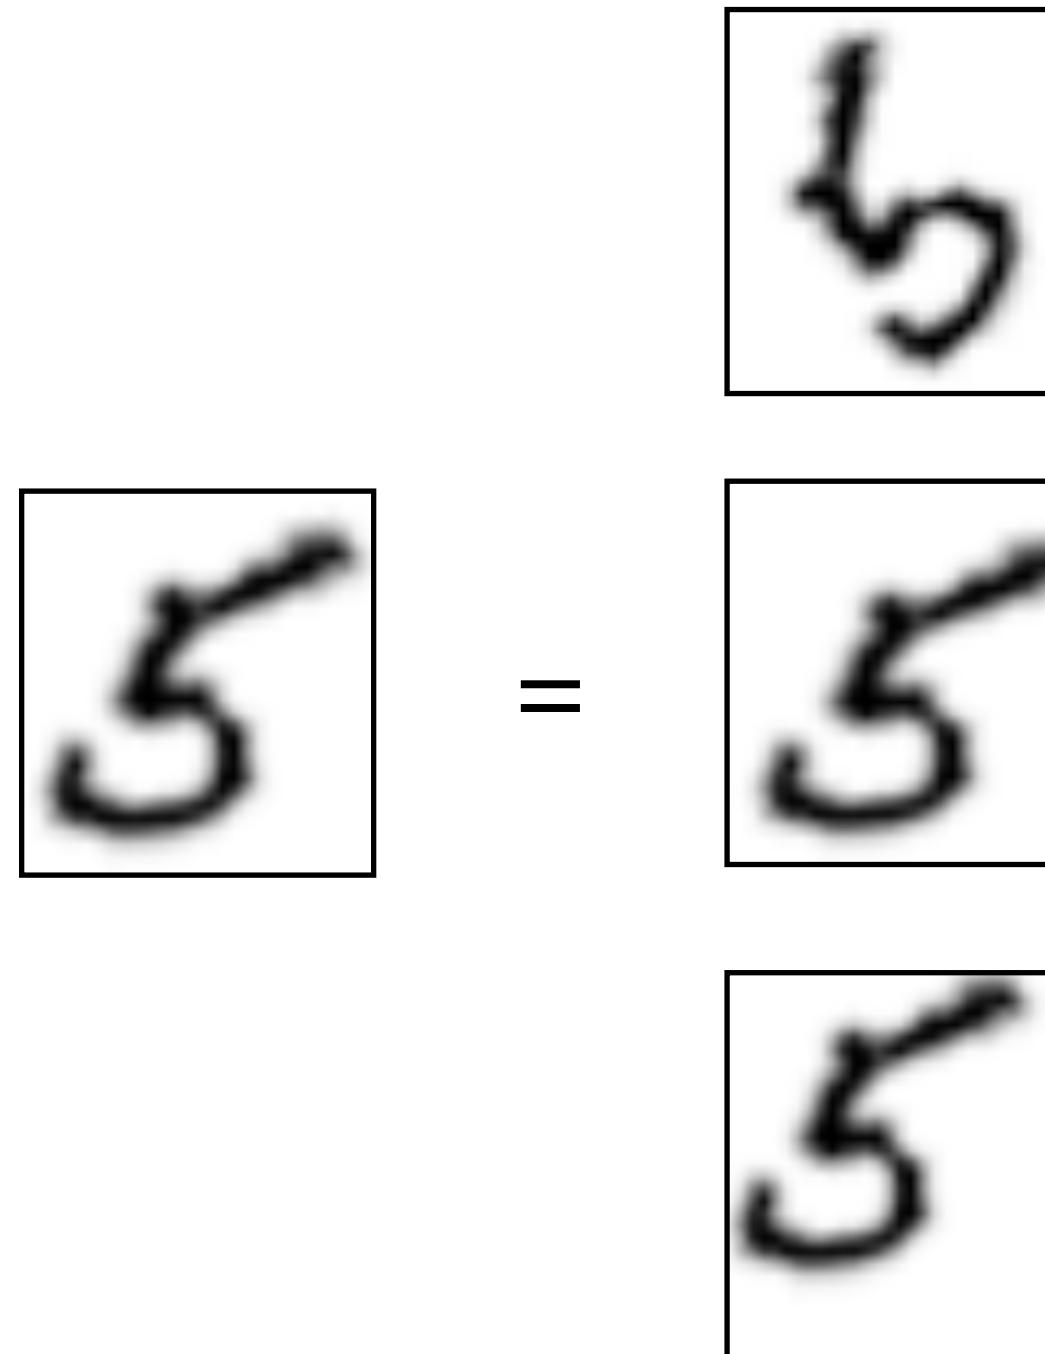
Convolutional neural nets
constrain multiple
parameters to be **equal**



(Images: Goodfellow 2016)

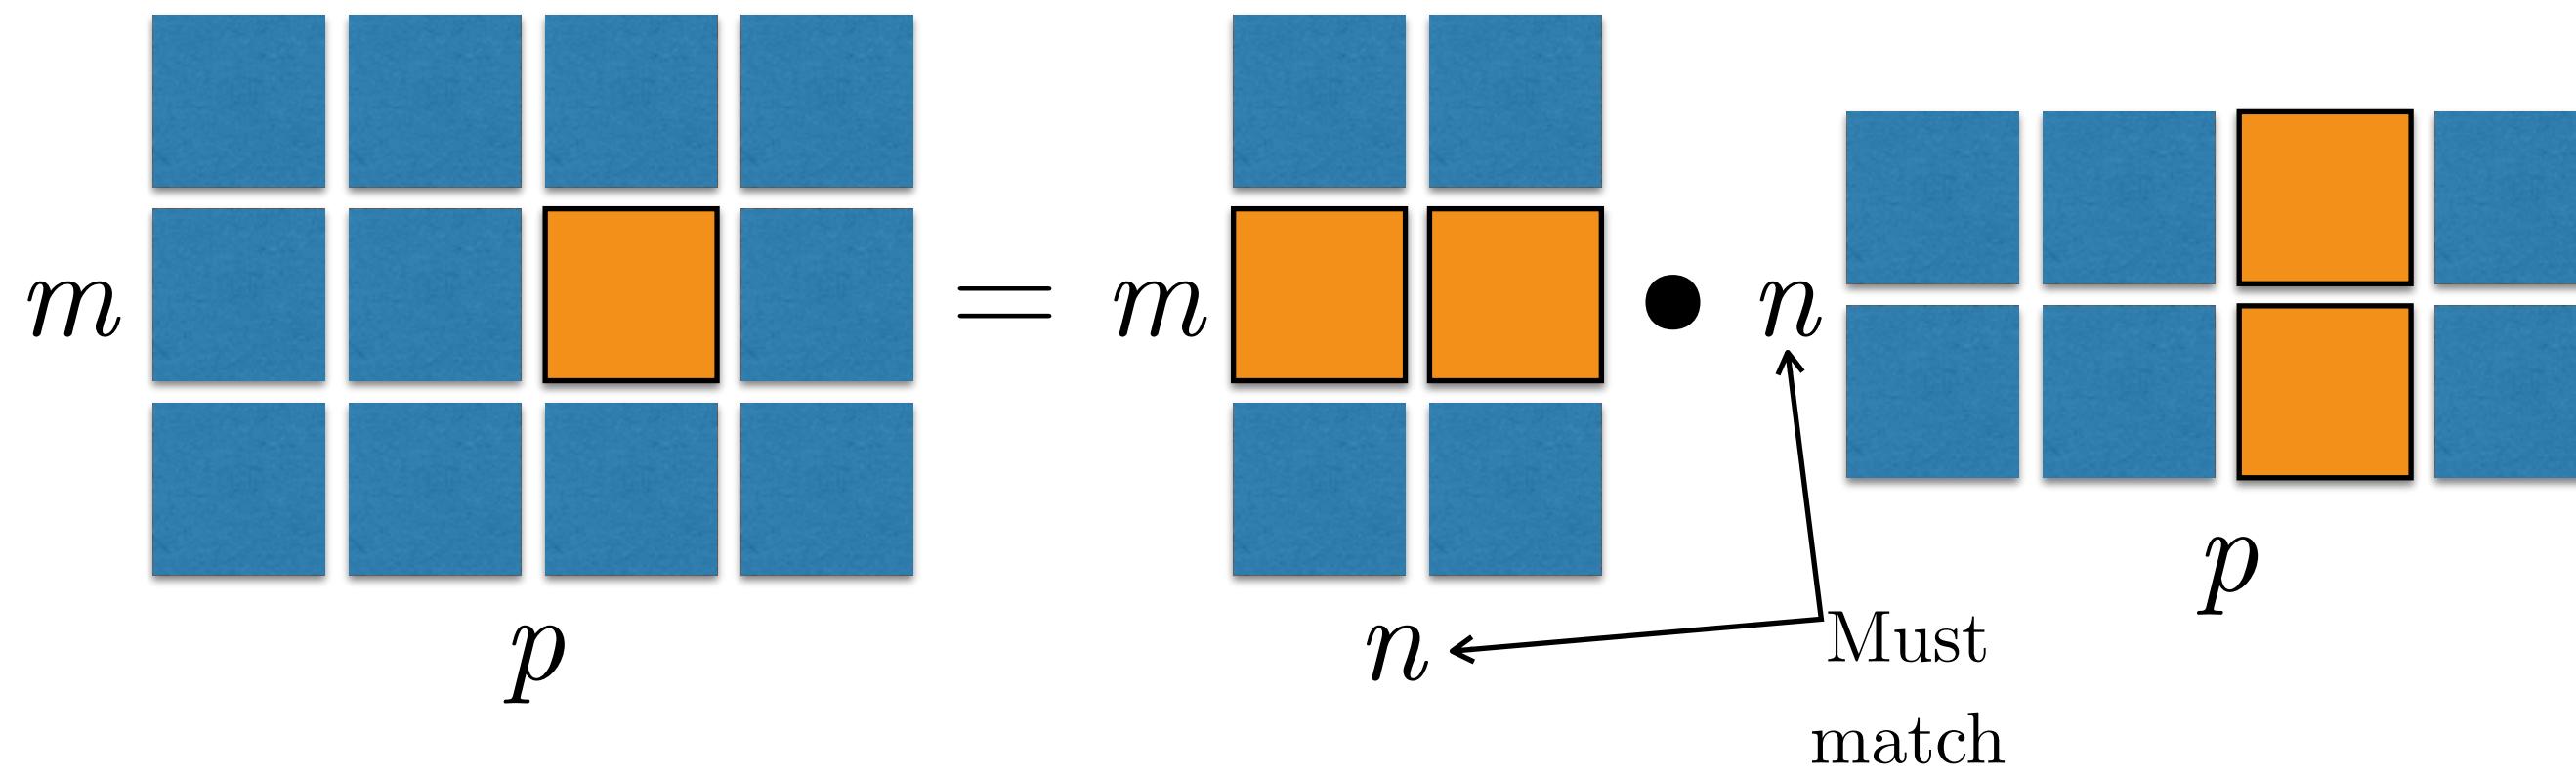
Equivariant Representations

- We want to be able to recognize transformed versions of inputs we have seen before:
 - Translation (moved)
 - Rotation
- Without having been **trained** on all transformed versions

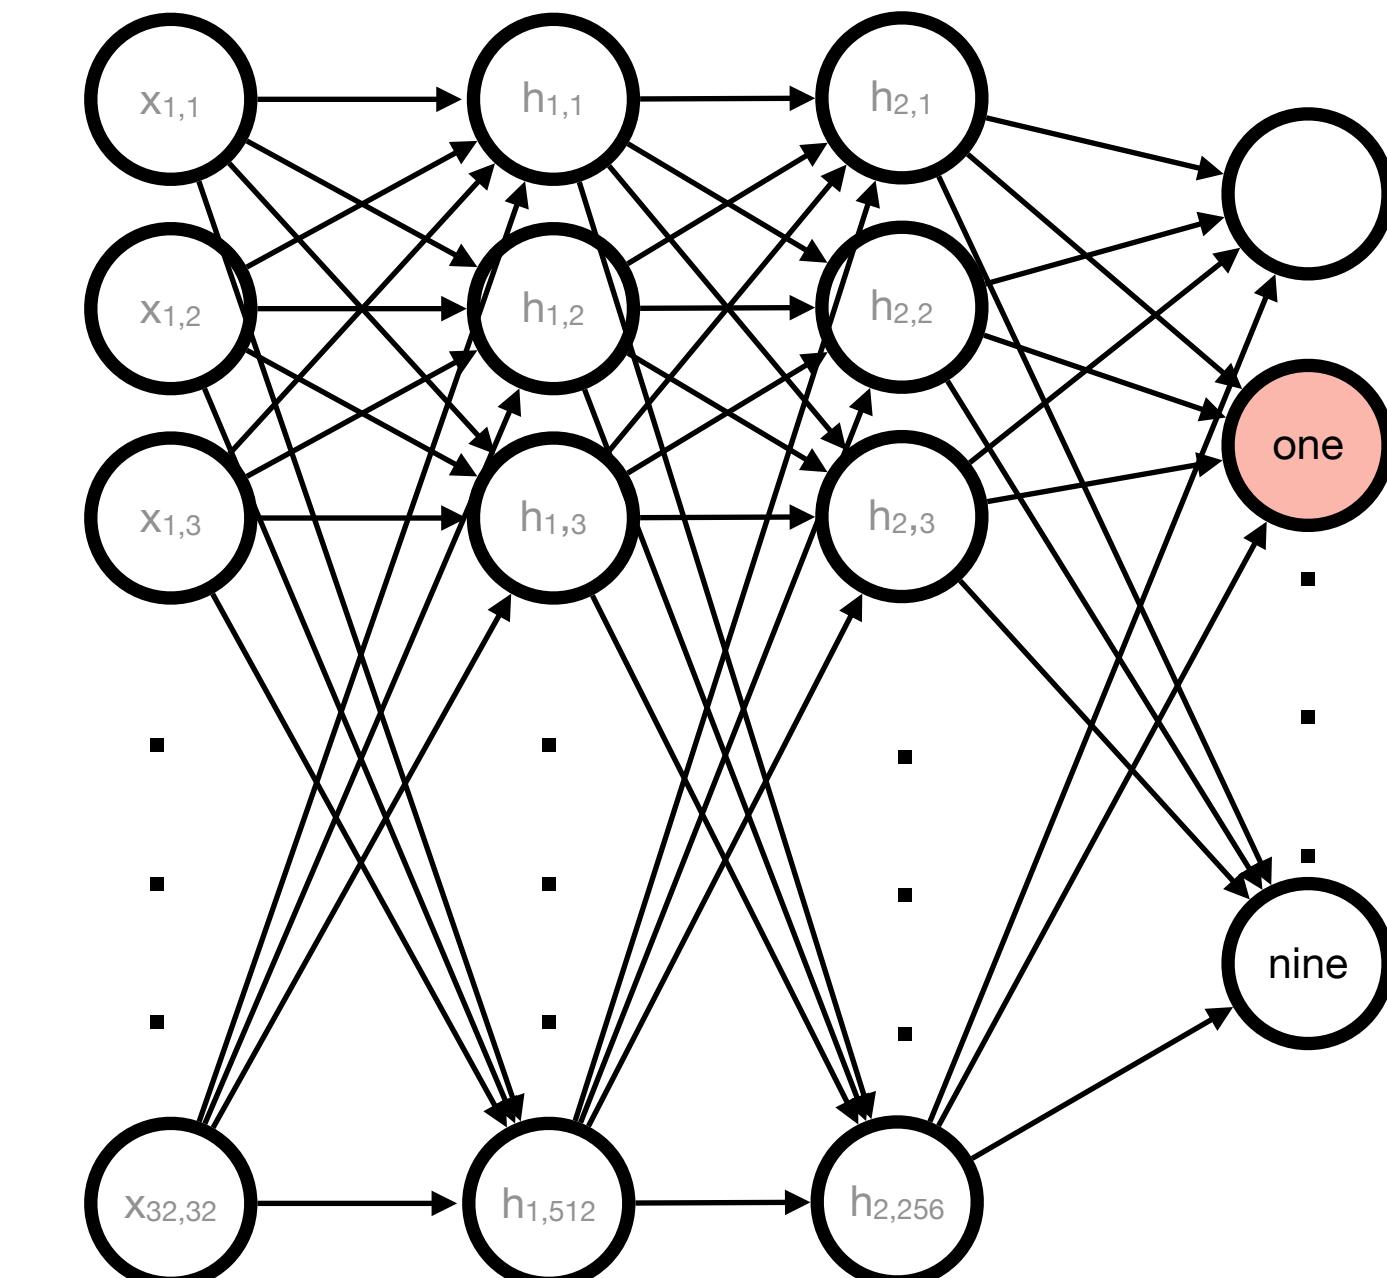


Operation: Matrix Product

Recall that we can represent the **activations** in a neural network by a **matrix product**



(Image: Goodfellow 2016)



$$\mathbf{h}_1 = g_h(W^{(1)}\mathbf{x} + \mathbf{b}^{(1)})$$

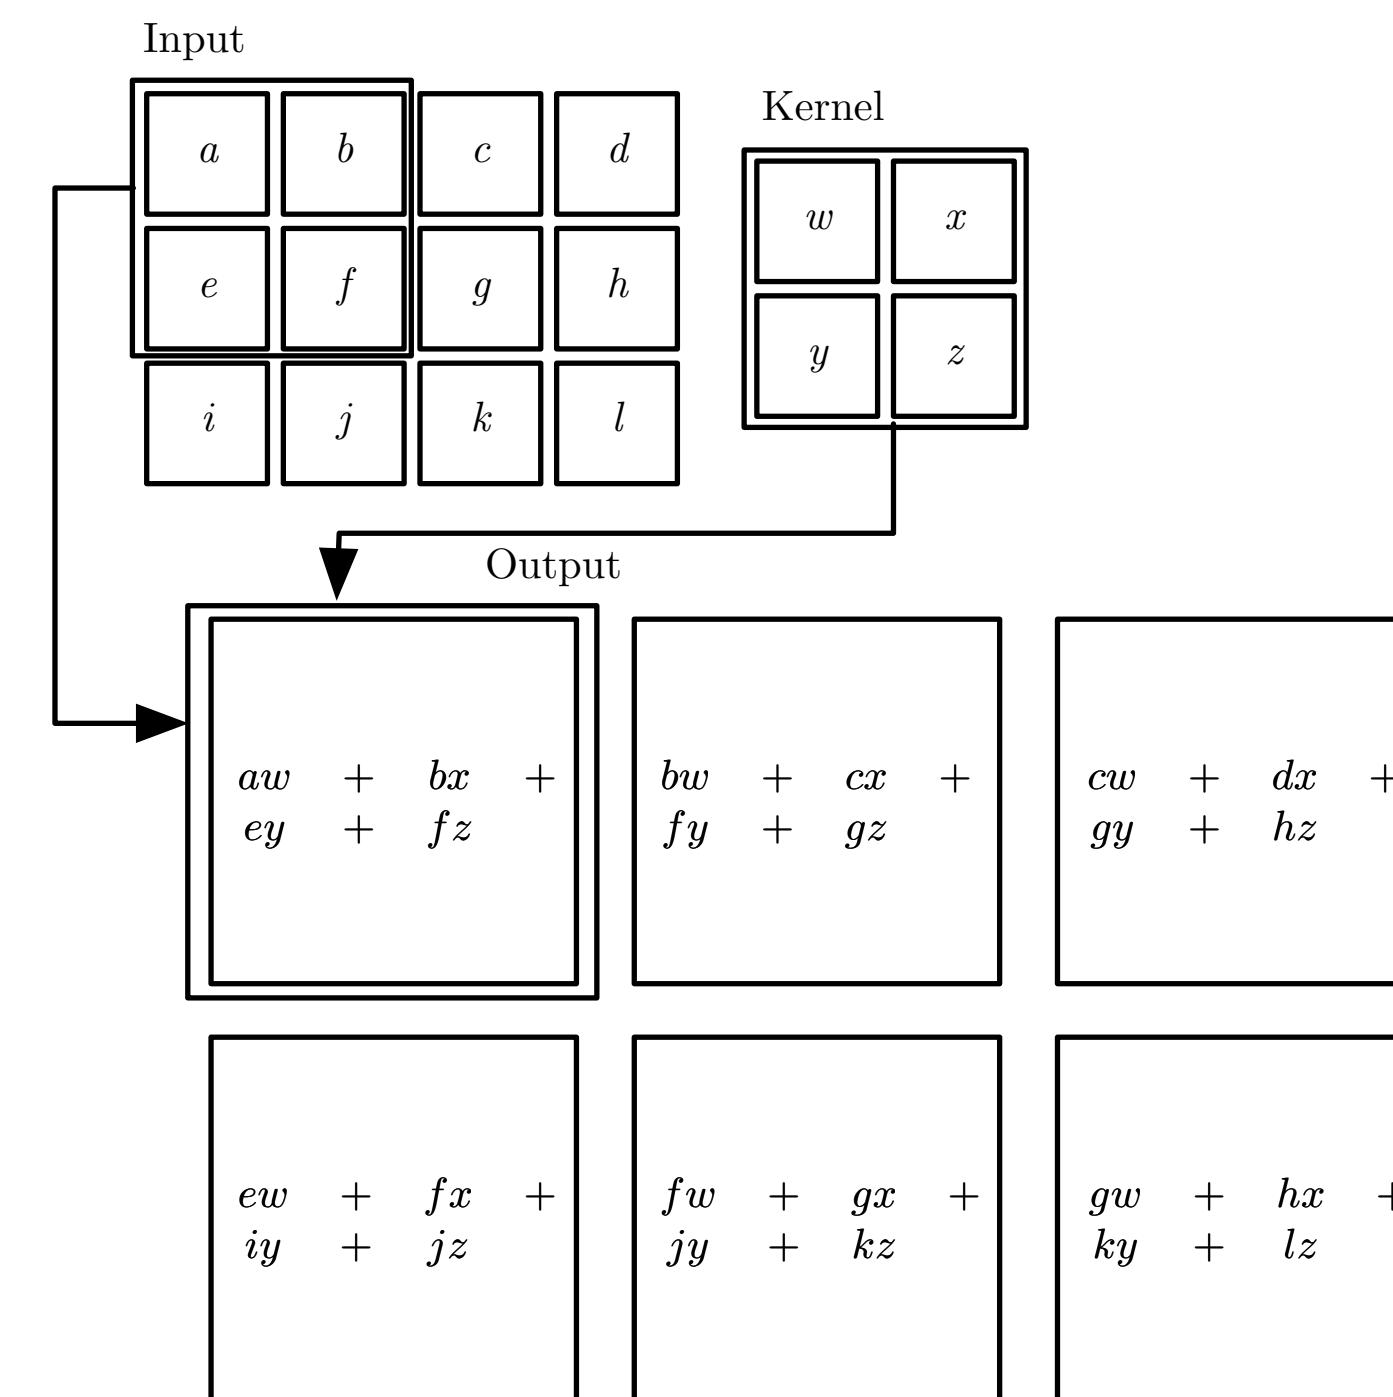
$$\mathbf{h}_2 = g_h(W^{(2)}\mathbf{h}_1 + \mathbf{b}^{(2)})$$

$$\mathbf{y} = g_y(W^{(3)}\mathbf{h}_2 + \mathbf{b}^{(3)})$$

Operation: 2D Convolution

Convolution scans a small block of weights (called the **kernel**) over the elements of the inputs, taking **weighted averages**

- Note that input and output dimensions **need not match**
- **Same weights** used for very many combinations



(Image: Goodfellow 2016)

Replace Matrix Multiplication by Convolution

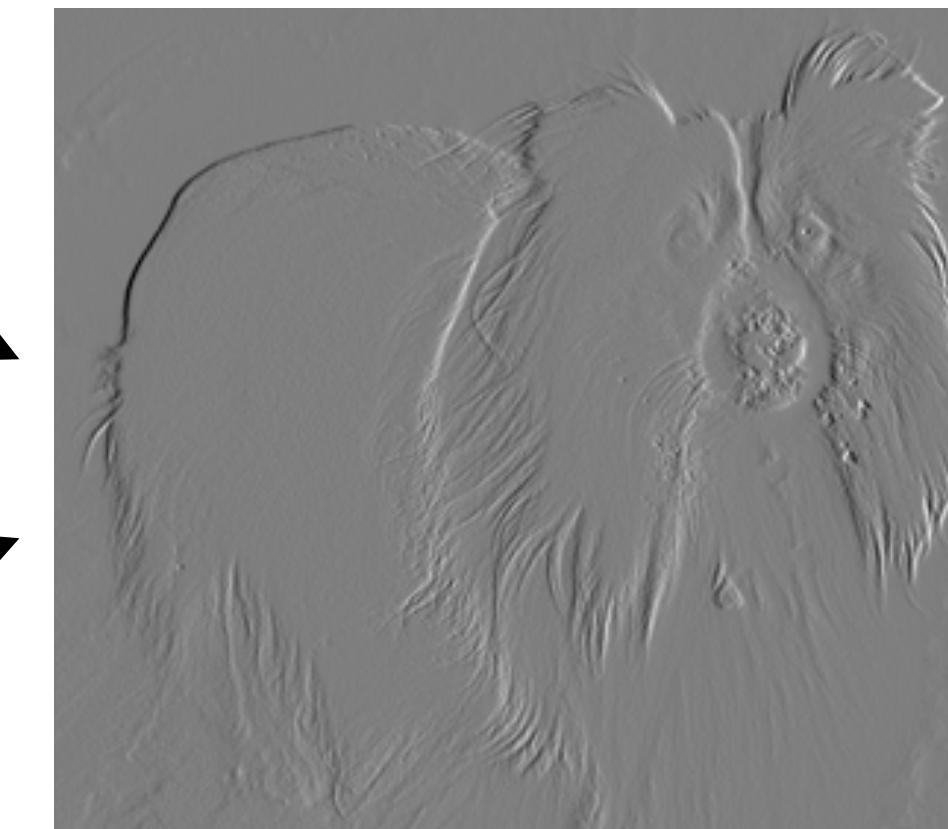
Main idea: Replace matrix multiplications with convolutions

- **Sparsity:** Inputs only combined with neighbours
- **Parameter sharing:** Same kernel used for entire input

Example: Edge Detection



Input



Output

1	-1
---	----

Kernel

(Image: Goodfellow 2016)

Efficiency of Convolution

Input size: 320 by 280

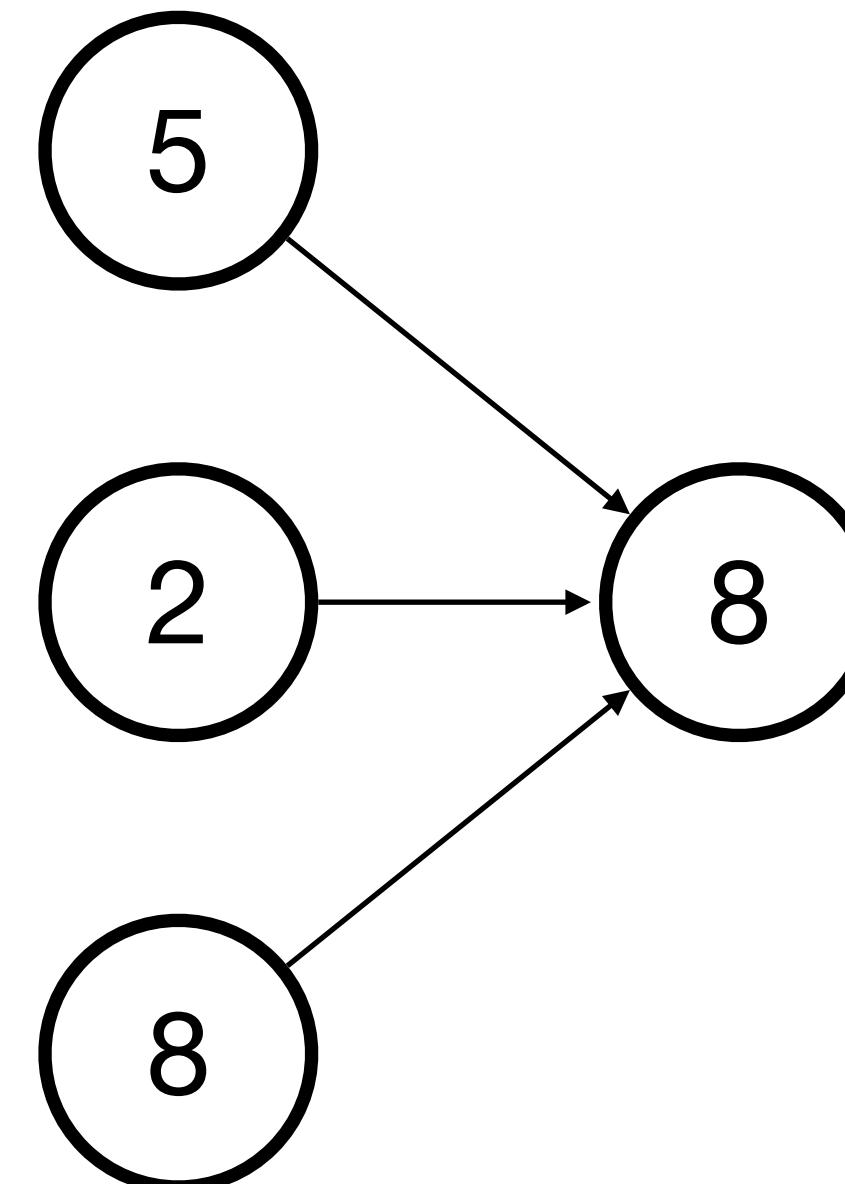
Kernel size: 2 by 1

Output size: 319 by 280

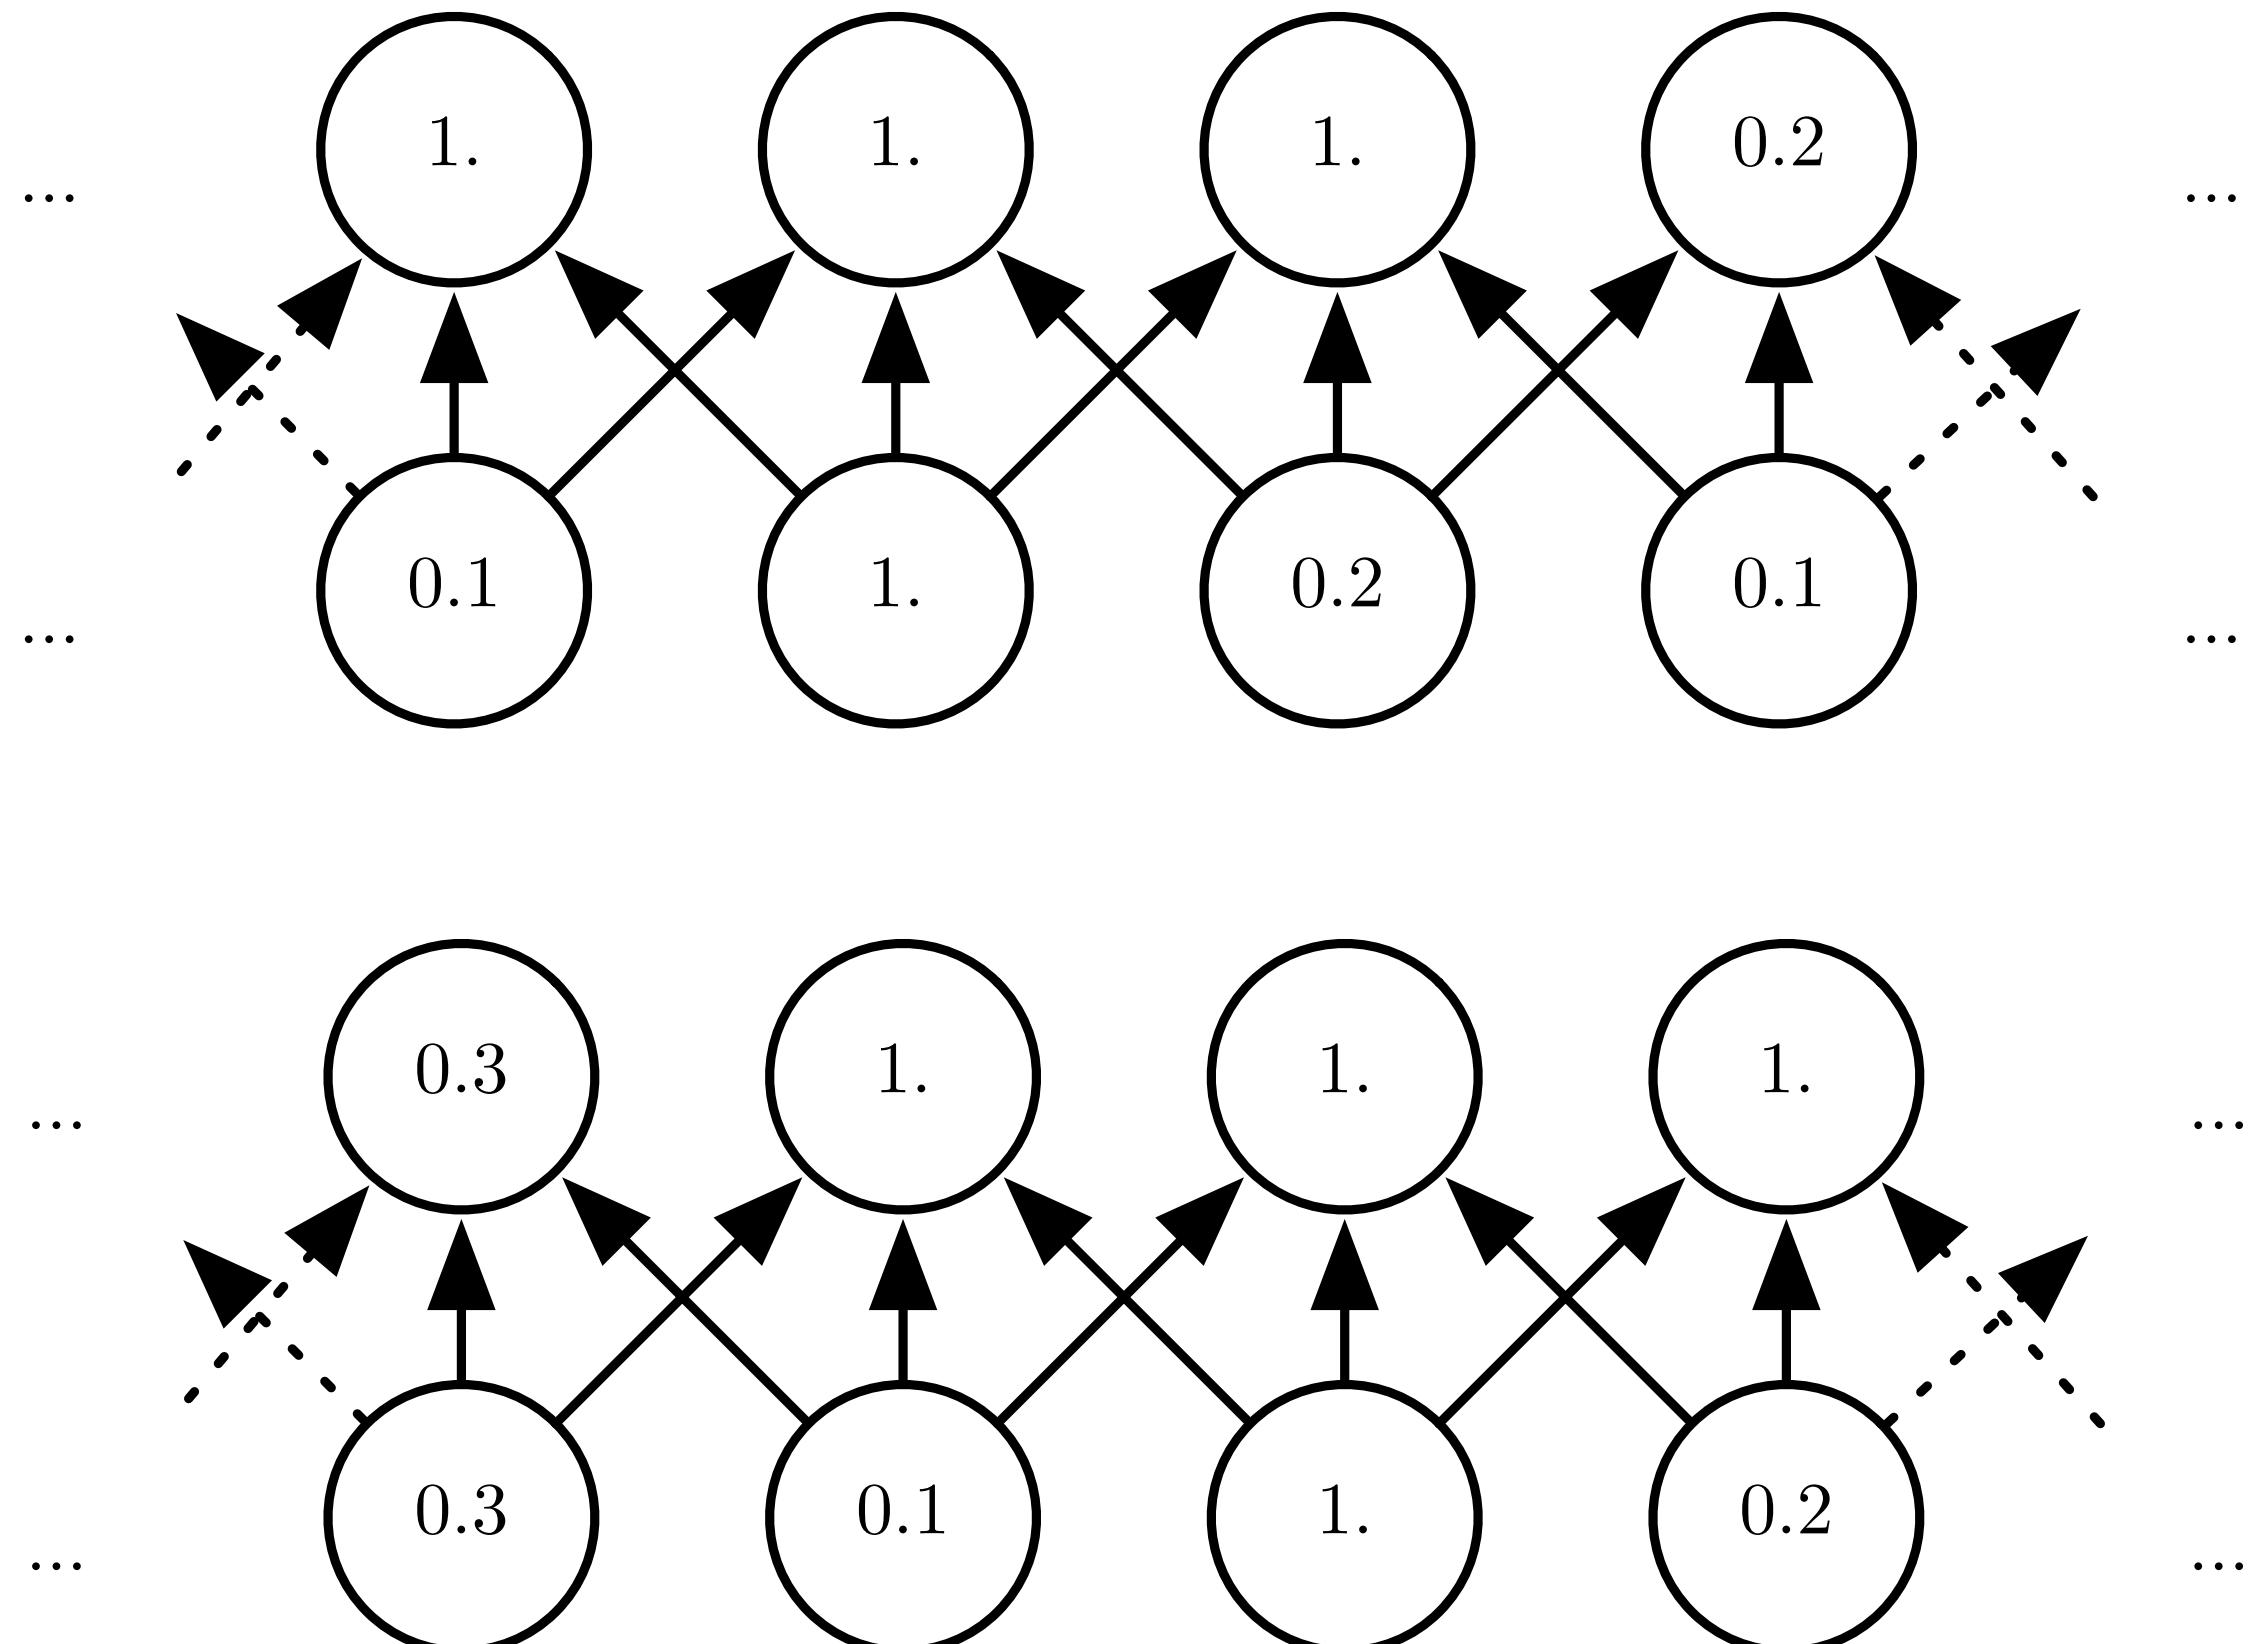
	Dense matrix	Sparse matrix	Convolution
Stored floats	$319 * 280 * 320 * 280$ $> 8e9$	$2 * 319 * 280 =$ 178,640	2
Float muls or adds	$> 16e9$	Same as convolution (267,960)	$319 * 280 * 3 =$ 267,960

Operation: Pooling

- Pooling **summarizes** its inputs into a single value, e.g.,
 - max
 - average
- Max-pooling is **parameter-free** (no bias or edge weights)

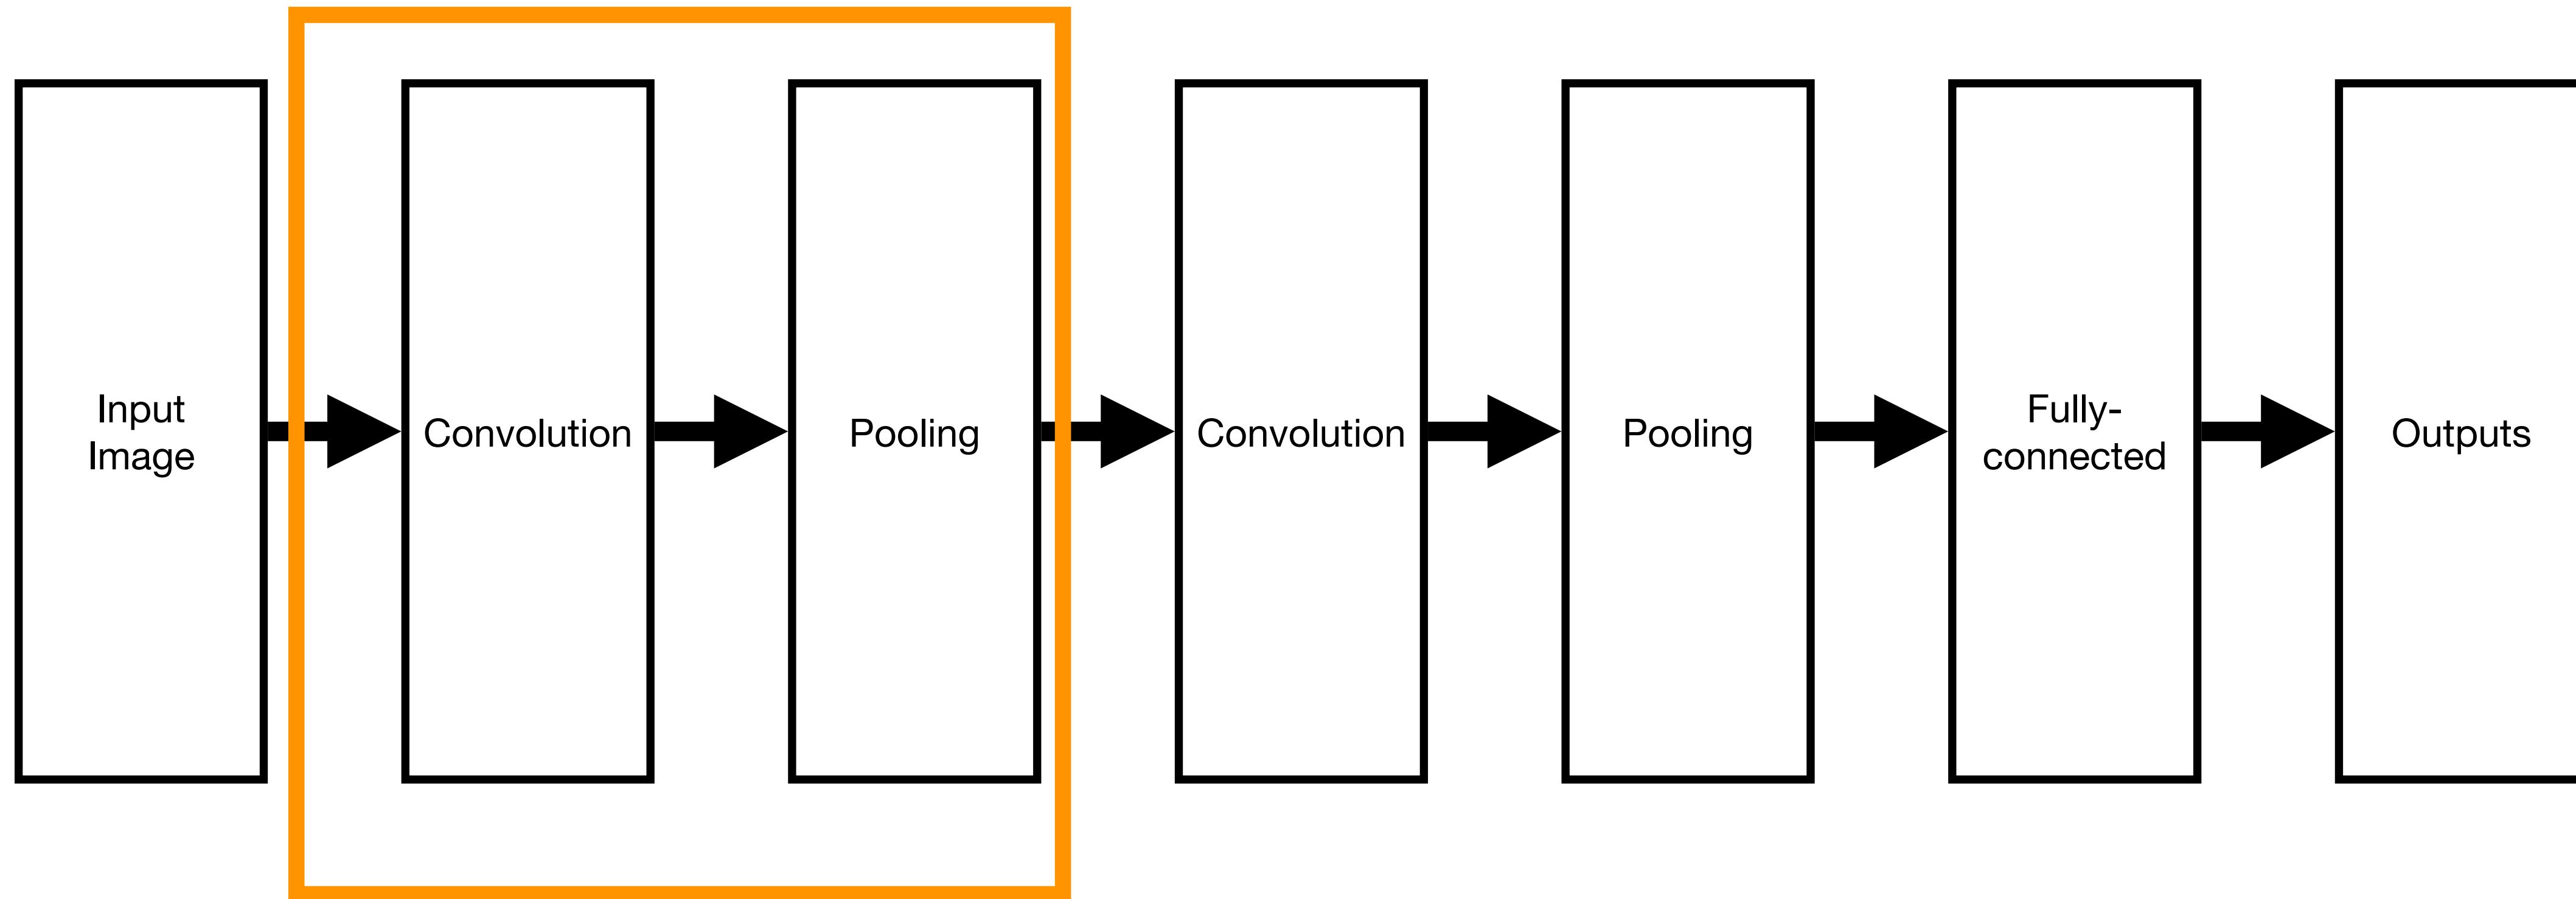


Example: Translation Invariance



(Goodfellow 2016)

Typical Architecture



Often convolution-then-pooling is
collectively referred to as a
"convolution layer"

Summary

- Classifying images with a standard feedforward network requires vast quantities of **parameters** (and hence **data**)
- Convolutional networks add **pooling** and **convolution**
 - Sparse connectivity
 - Parameter sharing
 - Translation equivariance
- Fewer parameters means far more **efficient to train**