

# Recurrent Neural Networks

CMPUT 366: Intelligent Systems



P&M §10.0-10.2, 10.10

# Lecture Outline

1. Recap
2. Unfolding Computations
3. Recurrent Neural Networks
4. Long Short-Term Memory

# Recap: Convolutional Neural Networks

- Convolutional networks: Specialized architecture for **images**
- Number of **parameters** controlled by using **convolutions** and **pooling** operations instead of **dense connections**
- Fewer parameters means more **efficient to train**



(Images: Goodfellow 2016)

# Sequence Modelling


- For many tasks, especially involving language, we want to model the behaviour of **sequences**
- **Example:** Translation
  - The cat is on the carpet  $\implies$  Le chat est sur le tapis
- **Example:** Sentiment analysis
  - This pie is great  $\implies$  POSITIVE
  - This pie is okay, not great  $\implies$  NEUTRAL
  - This pie is not okay  $\implies$  NEGATIVE

# Sequential Inputs

**Question:** How should we **represent** sequential input to a neural network?

1. 1-hot vector for **each word**  
(Sequence must be a particular length)
2. 1-hot vector for **last few words**  
( $n$ -gram)
3. **Single vector** indicating each word that is present  
(bag of words)
4. Single vector summing the **semantic embeddings** of all the words

The cat is on the carpet




# Dynamical Systems

- A **dynamical system** is a system whose state at time  $t + 1$  depends on its state at time  $t$ :

$$\mathbf{s}^{(t)} = f(\mathbf{s}^{(t-1)}; \theta)$$


- An expression that depends on the same expression at an earlier time is **recurrent**.



# Unfolding Computations

- A recurrent expression can be converted to a non-recurrent expression by **unfolding**:


$$\begin{aligned}\mathbf{s}^{(3)} &= f(\mathbf{s}^{(2)}; \theta) \\ &= f(f(\mathbf{s}^{(1)}; \theta); \theta)\end{aligned}$$



(Image: Goodfellow 2016)

# External Signals

- Dynamical systems can also be driven by **external signals**:
$$\mathbf{s}^{(t)} = f(\mathbf{s}^{(t-1)}, \mathbf{x}^{(t)}; \theta)$$
- These systems can also be represented by non-recurrent, unfolded computations:




(Image: Goodfellow 2016)

# Recurrent Neural Networks

- Recurrent neural network: a specialized architecture for modelling **sequential data**
- Input presented **one element at a time**
- Parameter sharing by:
  - Treating the sequence as a system with **state**
  - Introducing hidden layers that **represent** state
  - Computing **state transitions** and **output** using **same functions** at each stage

$x^{(6)} =$  


# Recurrent Hidden Units: Sequence to Sequence



- **Input values  $\mathbf{x}$**  connected to **hidden state  $\mathbf{h}$**  by weights  $\mathbf{U}$
- Hidden state  $\mathbf{h}$  mapped to **output  $\mathbf{o}$**  by weights  $\mathbf{V}$
- Hidden state  $\mathbf{h}^{(t-1)}$  connected to hidden state  $\mathbf{h}^{(t)}$  by weights  $\mathbf{W}$
- Gradients computed by **back propagation through time**: from final loss all the way back to initial input.
- All hidden states computed must be **stored** for computing gradients

(Image: Goodfellow 2016)


# Recurrent Hidden Units: Sequence to Single Output



- Update state as inputs are provided
- Only compute a **single** output at the **end**
- **W, U** still shared at every stage
- Back propagation through time still requires **evaluating every state** in gradient computation

(Image: Goodfellow 2016)

# Encoder/Decoder Architecture for Sequence to Sequence




Can **combine approaches** for sequence-to-sequence:

1. Accept entire input to construct a single "**context**" output **C**
2. Construct new sequence using context **C** as **only input**


(Image: Goodfellow 2016)

# Recurrence through (only) Outputs



- Can have recurrence go from **output** (at  $t - 1$ ) to **hidden** (at  $t$ ) instead of hidden to hidden
- Less general (**why?**)
- **Question:** Why would we want to do this?

# Teacher Forcing




- Dependence on previous step is only on output, not hidden state
- **Loss gradient** depends only on a **single transition**
- Training can be **parallelized** (don't need to compute previous states to compute current state)

(Image: Goodfellow 2016)


# Long-Range Dependence

The **submarine**, which was the subject of a well known song by the Beatles, was **yellow**.



- Information sometimes needs to be **accumulated** for a long part of the sequence
- But **how long** an individual piece of information should be accumulated is **context-dependent**
- Often need to **accumulate** information in the state, and then **forget** it later

# Long Short-Term Memory



- LSTM networks replace regular hidden units with **cells**
- Input feature computed with regular neuron
- Feature **accumulated** into state only if **input gate** allows it
- State **decays** according to value of **forget gate**
- Output can be **shut off** by the **output gate**

# Summary

- Naïvely representing **sequential inputs** for a neural network requires infeasibly many input nodes (and hence **parameters**)
- Recurrent neural networks are a **specialized architecture** for handling sequential inputs
  - **State** accumulates across input elements
  - Each stage computed from **previous stage** using **same parameters**
- **Long short-term memory** (LSTM) cells allow **context-dependent** accumulation and forgetting